

JACOB PFAU

✉ jacob.pfau@gmail.com 🔗 scholar.google.com/citations?user=r1IMgMAAAAJ 📄 jacobpfau

EDUCATION

New York University, Bowman Lab PhD Student NYU Alignment Research Group, Center for Data Science Current work includes: Red-teaming and latent adversarial training for LMs, Filler tokens and chain-of-thought in LMs	Fall 2022 - Present
University of Edinburgh MSc Mind, Language and Embodied Cognition (Philosophy Department). First class honors Dissertation: 'When Decision Theories and AI Bluff in Games with Perfect Information'	Sept. 2021 - Aug. 2022
Ecole Polytechnique (Université Paris-Saclay), Paris M1, Masters Computer Science, Data Science track. GPA 3.91	Sept. 2018 - Aug. 2019
Amherst College BA Mathematics & 5 College Logic Certificate. GPA (Math/CS/Logic) 3.73, Overall GPA 3.64 ETH Zurich (Math) Visiting student Jan-Aug 2016. GPA 5.34/6.0	Sept. 2013 - May 2017

EMPLOYMENT

Anthropic, Contractor Collaborative project writing a report on 'Responsible Scaling for AI Moral Patienthood' assessing the implications of philosophical work for LLMs.	Dec. 2023 - Present
AI Safety Camp, Mentor Mentored four researchers transitioning into AI safety. Developed a benchmark for evaluating LM self-consistency on an arithmetic task	Mar. 2023 - July 2023
University of California, San Francisco. Keiser Lab, Bold and Basic Fellow Developed interpretability methods and benchmarking datasets for clinical imaging. Mentored rotating students	Mar. 2019 - Aug. 2020

PUBLICATIONS

Eliciting Language Model Behaviors using Reverse Language Models. http://tinyurl.com/r1mpdf , Pfau J, Infanger A, Sheshadri A, Panda A, Huebner C, Michael J. Spotlight SoLaR at NeurIPS (2023)
Open problems and fundamental limitations of reinforcement learning from human feedback, Casper S, ..., Pfau J, et al. TMLR (2023)
Self-Consistency of Large Language Models under Ambiguity, Bartsch H,... Pfau J. BlackboxNLP (2023)
Objective Robustness in Deep Reinforcement Learning, Koch J, Langosco L, Pfau J. Le J, Sharkey L. ICML (2022)
Robust Semantic Interpretability: Revisiting Concept Activation Vectors, Pfau J, Young A, Wei J, Wei M, Keiser M. ICML Workshop on Human Interpretability (2020)
Stress Testing Reveals Gaps in Clinic Readiness of Image-Based Diagnostic AI Models, Young A, Fernandez K, Pfau J, et al. npj Digital Medicine (2020)
Artificial Intelligence in Dermatology: A Primer Young A, ... Pfau J, et al. Journal of Investigative Dermatology (2020)
Artificial Intelligence in Teledermatology, Xiong M, Pfau J, et al. Curr Derm Rep (2019)
Global Saliency: Aggregating Saliency Maps to Assess Dataset Artefact Bias, Pfau J, Young A, Wei M, Keiser M. NeurIPS ML4Health (2019)

FELLOWSHIPS AND AWARDS

Bold and Basic Fellowship, University of California, San Francisco Awarded yearlong research fellowship for applying CNNs to bridge skin cancer image data and vectorized genomic data	Sept. 2019 - Aug. 2020
Research Internship Prize, Ecole Polytechnique, Paris (for work conducted at UCSF)	Oct. 2019
Labex DigiCosme Fellowship, Labex DigiCosme, Paris	Sept. 2018 - Mar. 2019

SKILLS AND MISCELLANEOUS

LANGUAGES: English (native), German (fluent), French (conversant), Mandarin (basic)